

Chapter 2: Intelligent Agents





Outline

- Last class, introduced AI and rational agent
- Today's class, focus on intelligent agents
 - Agent and environments
 - Video lecture available
 - Nature of environments influences agent design
 - Video lecture available
 - Basic "skeleton" agent designs
 - Review by yourself




Outline

- Last class, introduced AI and rational agent
- Today's class, focus on intelligent agents
 - **Agent and environments**
 - Nature of environments influences agent design
 - Basic “skeleton” agent designs


Agents

- An **agent** is anything that can be viewed as **perceiving** its **environment** through **sensors** and **acting** upon that environment through **actuators**
- Examples:
 - Human agent
 - Robotic agent
 - Software agent



Terminologies

- **Percept**: the agent's perceptual inputs
- **Percept sequence**: the complete history of everything the agent has perceived
- **Agent function** maps any given percept sequence to an action $[f: p^* \rightarrow A]$
- The **agent program** runs on the physical architecture to produce f
- **Agent = architecture + program**






Questions


- Can there be more than one agent program that implements a given agent function?
- Given a fixed machine architecture, does each agent program implement exactly one agent function?


Vacuum-Cleaner World



- **Percepts:** location and contents, e.g., [A, dirty]
- **Actions:** Left, Right, Suck, NoOp

A Simple Agent Function

Percept sequence	Action
[A, Clean]	
[A, Dirty]	
[B, Clean]	
[B, Dirty]	
[A, Clean], [A, Clean]	
[A, Clean], [A, Dirty]	
...	
[A, Clean], [A, Clean], [A, Clean]	
[A, Clean], [A, Clean], [A, Dirty]	
...	





Rationality

- An agent should "do the right thing", based on what it can perceive and the actions it can perform. The right action is the one that will cause the agent to be most successful
- **Performance measure:** An objective criterion for success of an agent's behavior
- Back to the vacuum-cleaner example
 - Amount of dirt cleaned within certain time
 - +1 credit for each clean square per unit time
- General rule: measure what one wants rather than how one thinks the agent should behave



Rational Agent

- Definition:

- For each possible percept sequence, a rational agent should select an action that is expected to maximize its performance measure, given the evidence provided by the percept sequence and whatever built-in knowledge the agent has.



Rational Agent

- Definition:

- For each possible **percept sequence**, a rational agent should select an **action** that is expected to maximize its **performance measure**, given the evidence provided by the percept sequence and whatever built-in **knowledge** the agent has.



Vacuum-Cleaner Example

- A simple agent that cleans a square if it is dirty and moves to the other square if not
- Is it rational?
- **Assumption:**
 - performance measure: 1 point for each clean square at each time step
 - environment is known a priori
 - actions = {left, right, suck, no-op}
 - agent is able to perceive the location and dirt in that location
- Given different assumption, it might not be rational anymore



Omniscience, Learning and Autonomy

- Distinction between rationality and **omniscience**
 - expected performance vs. actual performance
- Agents can perform actions in order to modify future percepts so as to obtain useful information (**information gathering, exploration**)
- An agent can also **learn** from what it perceives
- An agent is **autonomous** if its behavior is determined by its own experience (with ability to learn and adapt)



In-Class Exercise #2.1

- Given the assumption on slide 12. Describe a rational agent function for the modified performance measure that deducts one point for each movement. Does the agent program require internal state?
- Discuss possible agent designs for the cases in which clean squares can become dirty and the geography of the environment is unknown.



Outline

- Last class, introduced AI and rational agent
- Today's class, focus on intelligent agents
 - Agent and environments
 - Nature of environments influences agent design
 - Basic “skeleton” agent designs



PEAS

- Specifying the task environment is always the first step in designing agent
- **PEAS:**
 - **P**erformance, **E**nvironment, **A**ctuators, **S**ensors



Taxi Driver Example

Performance Measure	Environment	Actuators	Sensors
safe, fast, legal, comfortable trip, maximize profits	roads, other traffic, pedestrians, customers	steering, accelerator, brake, signal, horn, display	camera, sonar, speedometer, GPS, odometer, engine sensors, keyboard, accelerator

DARPA urban challenge 07:

<http://www.youtube.com/watch?v=SQFEmR50HAK>

Medical Diagnosis System

Performance Measure	Environment	Actuators	Sensors
healthy patient, minimize costs, lawsuits	patient, hospital, staff	display questions, tests, diagnosis, treatments, referrals	keyboard entry of symptoms, findings, patient's answers



Mushroom-Picking Robot

Performance Measure	Environment	Actuators	Sensors
Percentage of good mushrooms in correct bins	Conveyor belt with mushrooms, bins	Jointed arm and hand	camera, joint angle sensors





Properties of Task Environments

- **Fully observable** (vs. **partially observable**):
 - An agent's sensors give it access to the complete state of the environment at each point in time
- **Deterministic** (vs. **stochastic**):
 - next state of the env. determined by current state and the agent's action
 - If the environment is deterministic except for the actions of other agents, then the environment is **strategic**
- **Episodic** (vs. **sequential**):
 - Agent's experience is divided into atomic "episodes"
 - Choice of action in each episode depends only on the episode itself



Properties of Task Environments

- **Static** (vs. **dynamic**):
 - The environment is unchanged while an agent is deliberating
 - **Semidynamic** if the environment itself doesn't change with time but the agent's performance score does
- **Discrete** (vs. **continuous**):
 - A limited number of distinct, clearly defined percepts and actions
- **Single agent** (vs. **multiagent**):
 - An agent operating by itself in an environment
 - Competitive vs. cooperative



Examples

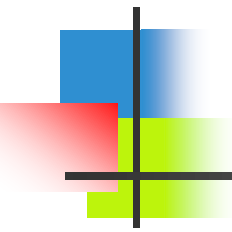
Task Environment	Oberservable	Deterministic	Episodic	Static	Discrete	Agents
<i>Crossword puzzle</i>	fully	deterministic	sequential	static	discrete	single
<i>Chess with a clock</i>	fully	strategic	sequential	semi	discrete	multi
<i>Taxi driver</i>	partially	stochastic	sequential	dynamic	conti.	multi
<i>mushroom-picking</i>	partially	stochastic	episodic	dynamic	conti.	single

- The environment type largely determines the agent design
- The real world is (of course) partially observable, stochastic, sequential, dynamic, continuous, multi-agent



In-Class Exercise #2.2

- Develop PEAS description for the following task environment:
 - Robot soccer player
 - Shopping for used AI books on the Internet
- Analyze the properties of the above environments



In-Class Exercise #3:

True/False Questions

- 1. An agent that senses only partial information about the state cannot be perfectly rational.
- 2. Suppose an agent selects its action uniformly at random from the set of possible actions. There exists a deterministic task environment in which this agent is rational.
- 3. It is possible for a given agent to be perfectly rational in two distinct task environments.
- 4. A perfectly rational poker-playing agent never loses.



Outline

- Last class, introduced AI and rational agent
- Today's class, focus on intelligent agents
 - Agent and environments
 - Nature of environments influences agent design
 - Basic "skeleton" agent designs



Agent = Architecture + Program

- The job of AI is to design the **agent program** that implements the **agent function** mapping percepts to actions
- Aim: find a way to implement the **rational agent function** concisely
- Same skeleton for agent program: it takes the **current percept** as input from the sensors and returns an action to the actuators



Agent Program vs. Agent Function

- Agent program takes the current percept as input
 - Nothing is available from the environment
- Agent function takes the entire percept history
 - To do this, remember all the percepts



Table-Driven Agent

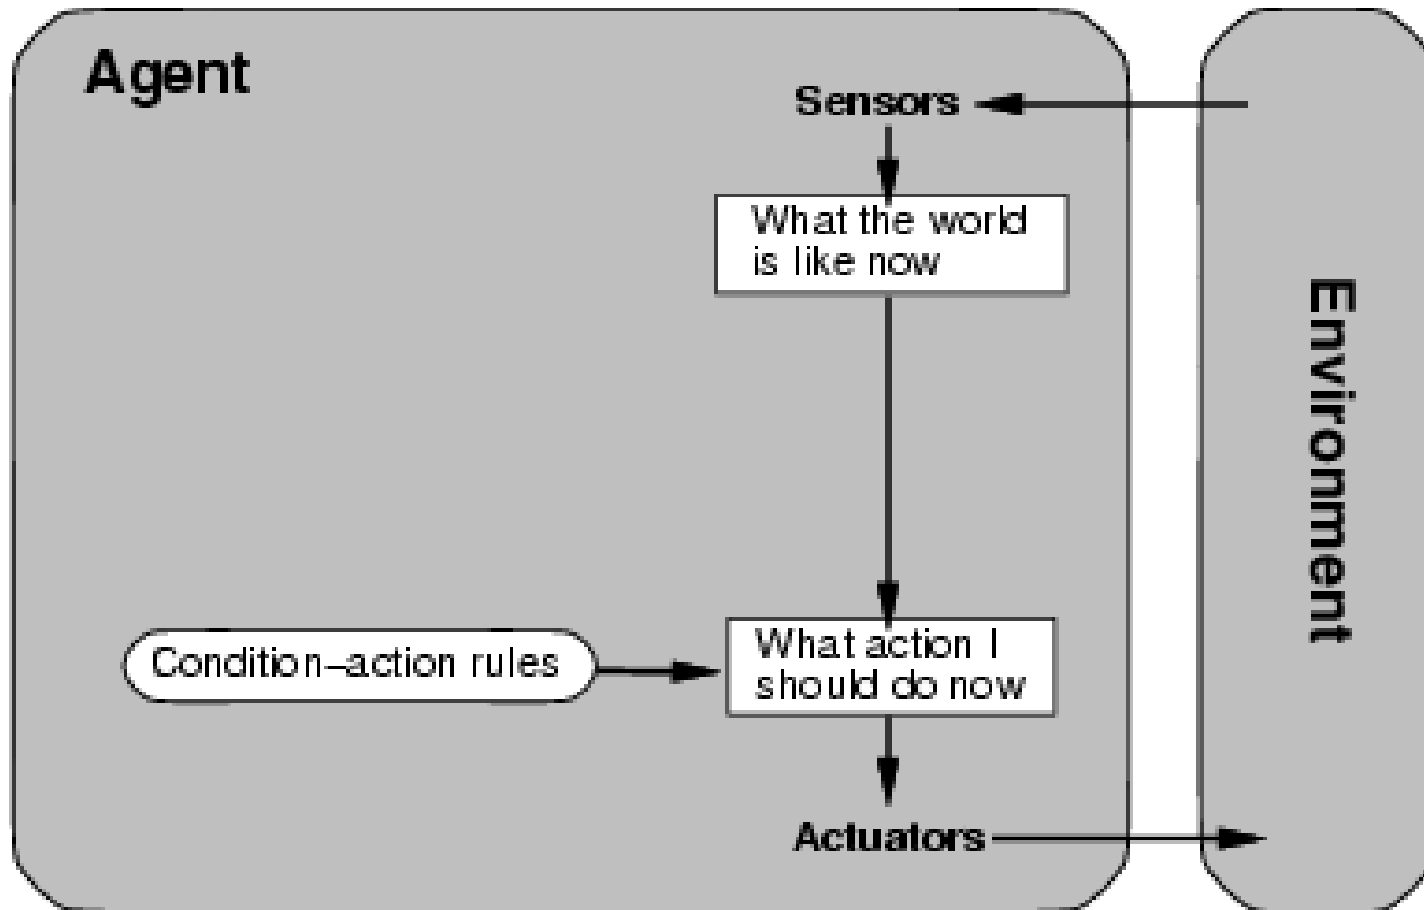
- Designer needs to construct a table that contains the appropriate action for every possible percept sequence
- Drawbacks?
 - huge table
 - take a long time to construct such a table
 - no autonomy
 - Even with learning, need a long time to learn the table entries



Five Basic Agent Types

- Arranged in order of increasing generality:
 - Simple reflex agents
 - Model-based reflex agents
 - Goal-based agents
 - Utility-based agents; and
 - Learning agents

Simple Reflex Agent



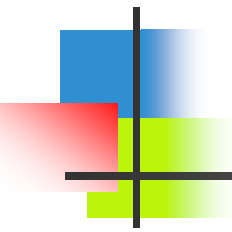


Pseudo-Code

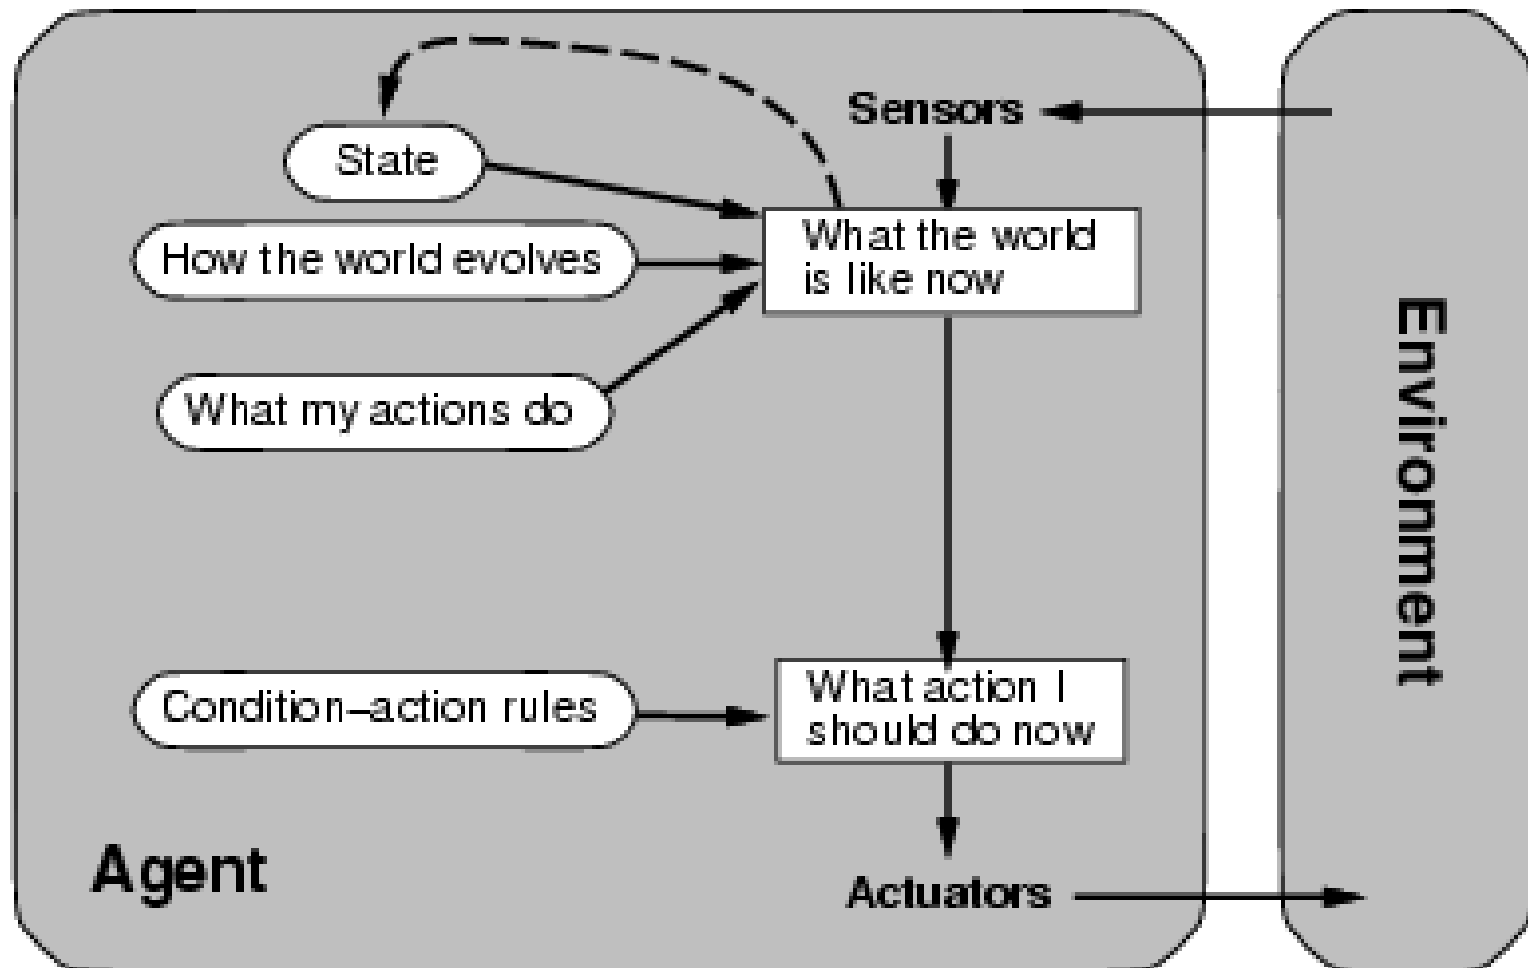
function SIMPLE-REFLEX-AGENT(*percept*) **returns** an action
 static: *rules*, a set of condition–action rules

state \leftarrow INTERPRET-INPUT(*percept*)
 rule \leftarrow RULE-MATCH(*state*, *rules*)
 action \leftarrow RULE-ACTION[*rule*]
 return *action*

Example: write a simple reflex agent for the vacuum cleaner example

- 
- **Infinite loops** are often unavoidable for simple reflex agent operating in partially observable environments
 - No location sensor
 - **Randomization** will help
 - A randomized simple reflex agent might outperform a deterministic simple reflex agent
 - Better way: keep track of the part of the world it can't see now
 - Maintain internal states

Model-Based Reflex Agent



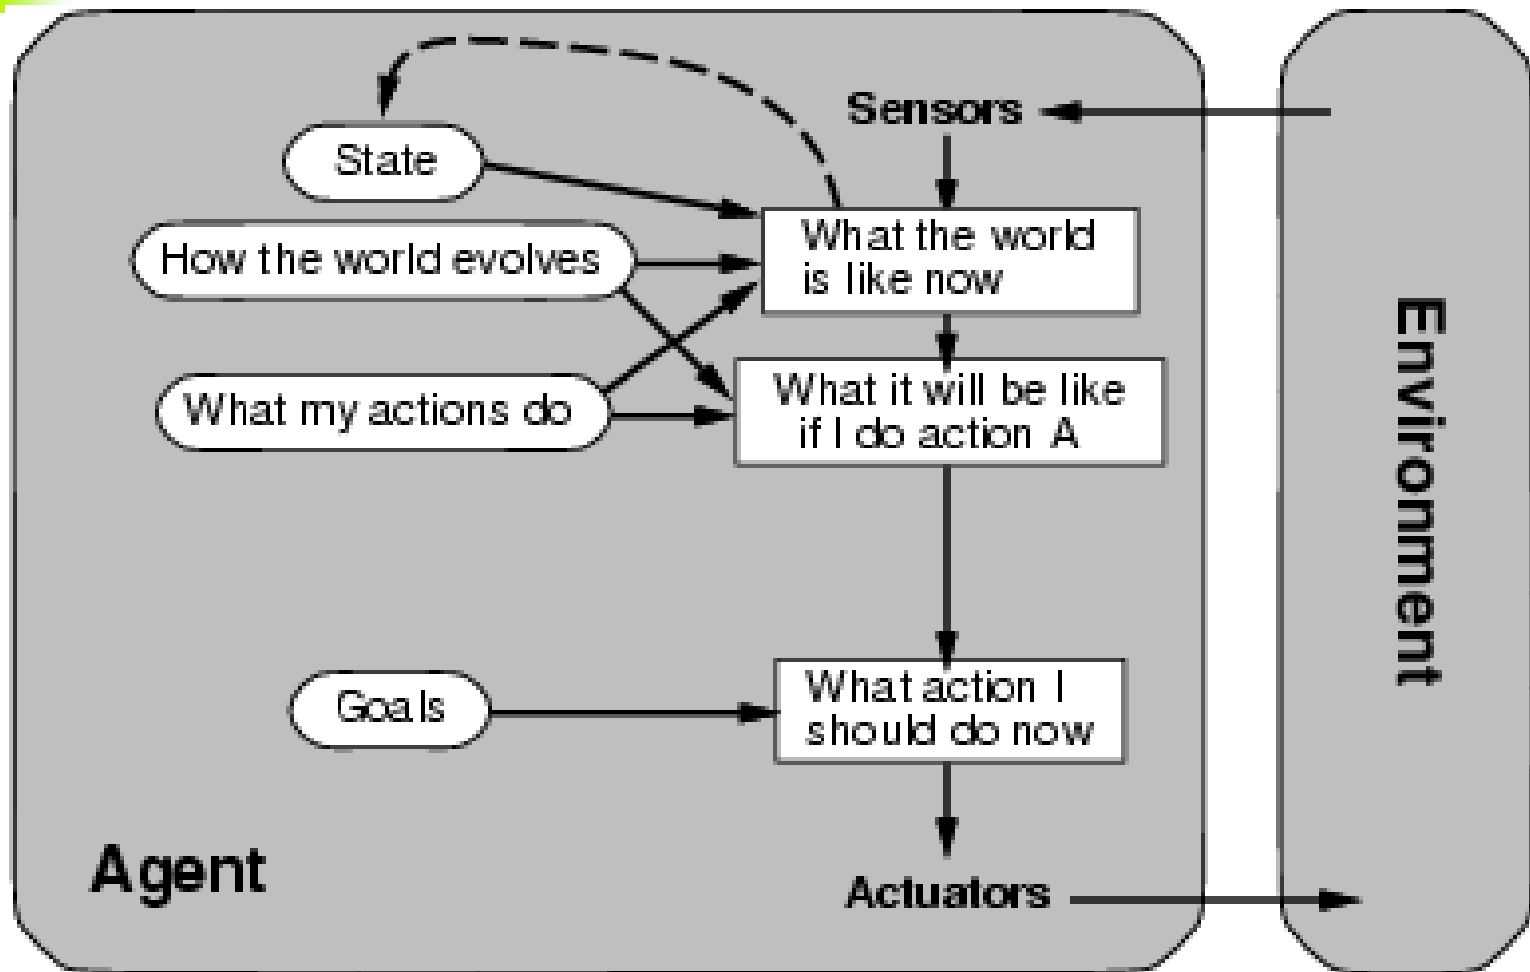


Pseudo-Code

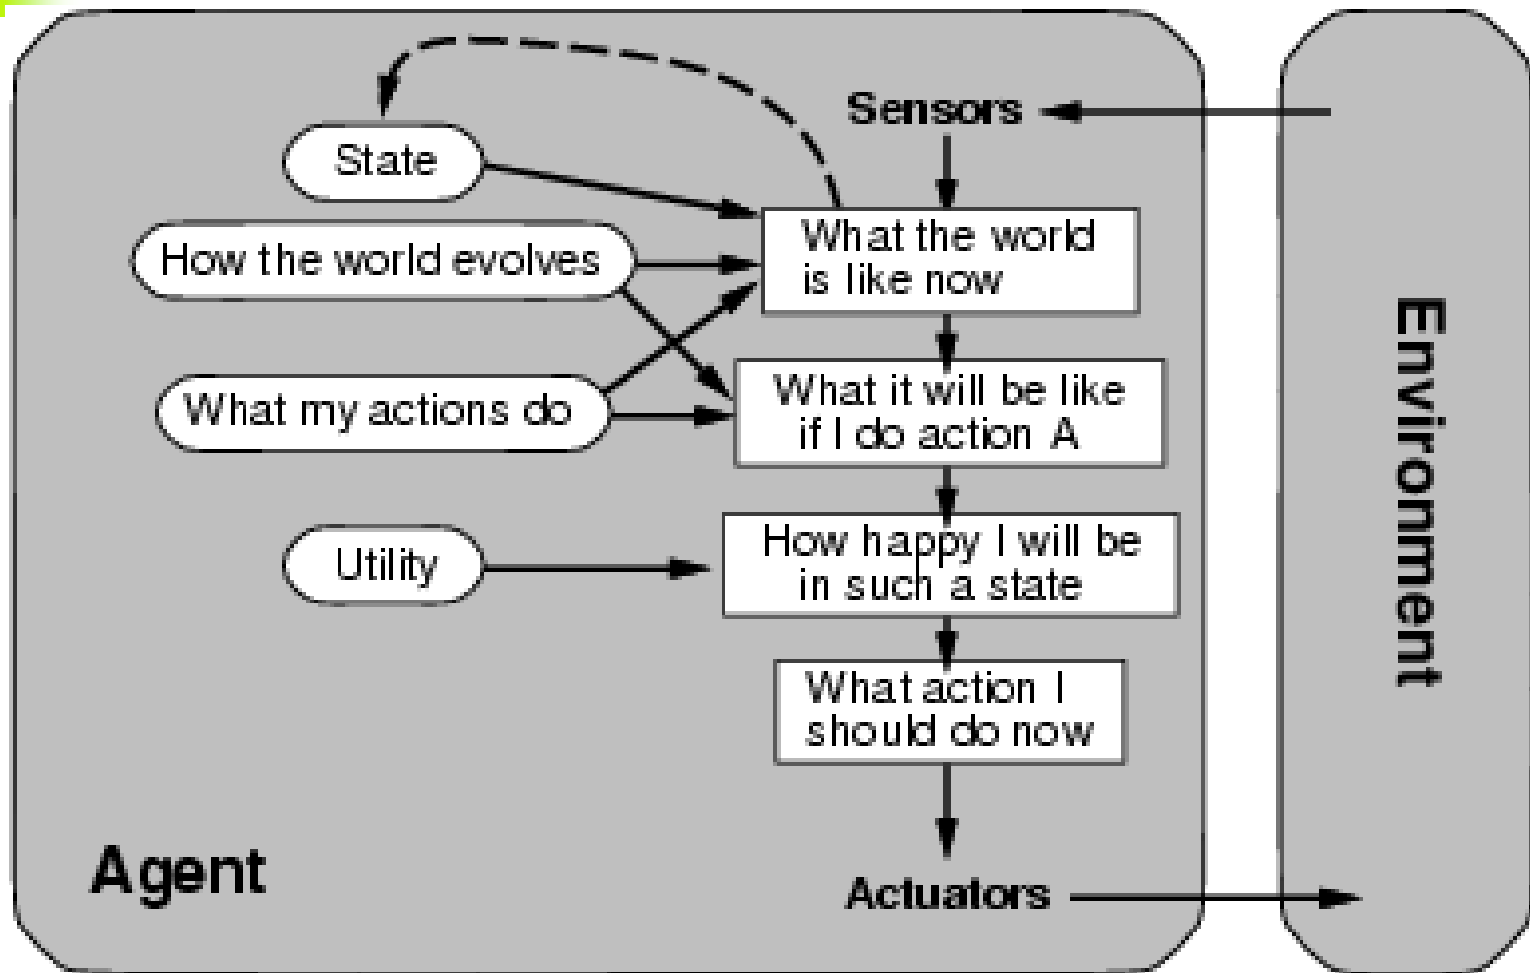
function REFLEX-AGENT-WITH-STATE(*percept*) **returns** an action
 static: *state*, a description of the current world state
 rules, a set of condition–action rules
 action, the most recent action, initially none

 state \leftarrow UPDATE-STATE(*state*, *action*, *percept*)
 rule \leftarrow RULE-MATCH(*state*, *rules*)
 action \leftarrow RULE-ACTION[*rule*]
 return *action*

Goal-Based Agent



Utility-Based Agent

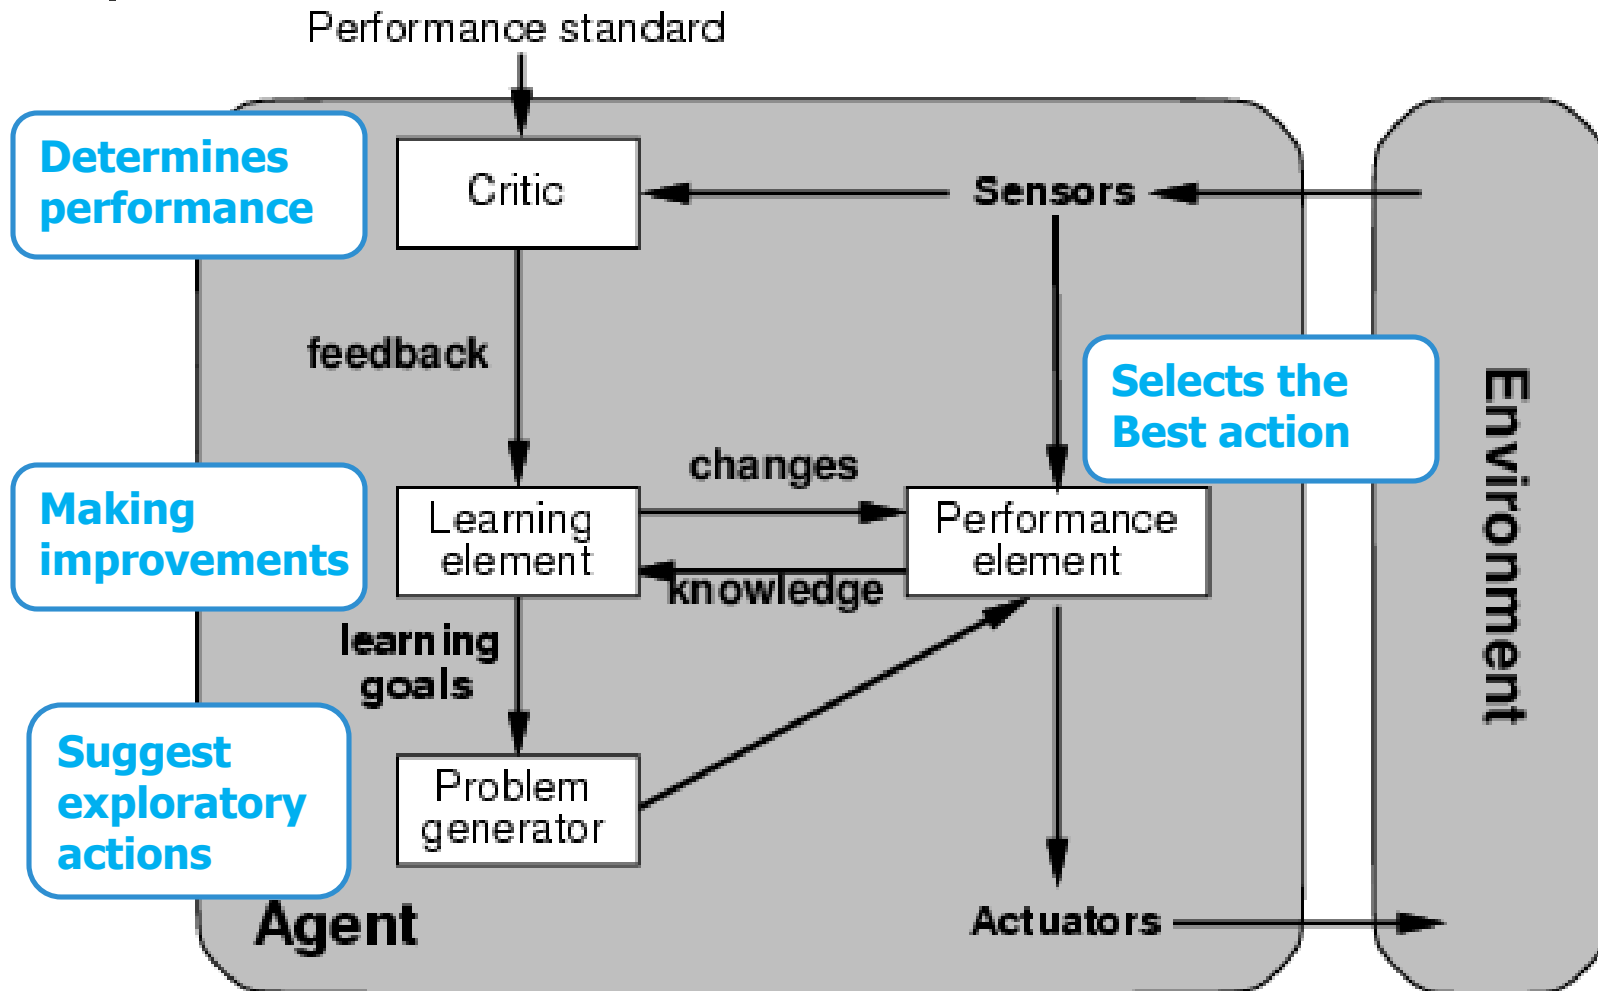




Utility Function

- Utility function maps a state or a sequence of states onto a real number \rightarrow degree of happiness
- Conflicting goals
 - Speed and safety
- Multiple goals

Learning Agent





Thinking Exercise

- Select a suitable agent design for:
 - Robot soccer player
 - Shopping for used AI books on the Internet



Summary

- Agent, agent function, agent program
 - Rational agent and its performance measure
 - PEAS
 - Five major agent program skeletons
-
- Next, solving problems by searching