

Chapter 3: Informed Search and Exploration

Dr. Daisy Tang



Informed Search

- ❑ Definition:
 - ❑ Use problem-specific knowledge beyond the definition of the problem itself
 - ❑ Can find solutions more efficiently
- ❑ Best-first search
 - ❑ Greedy best-first search
 - ❑ A^*
- ❑ Heuristics

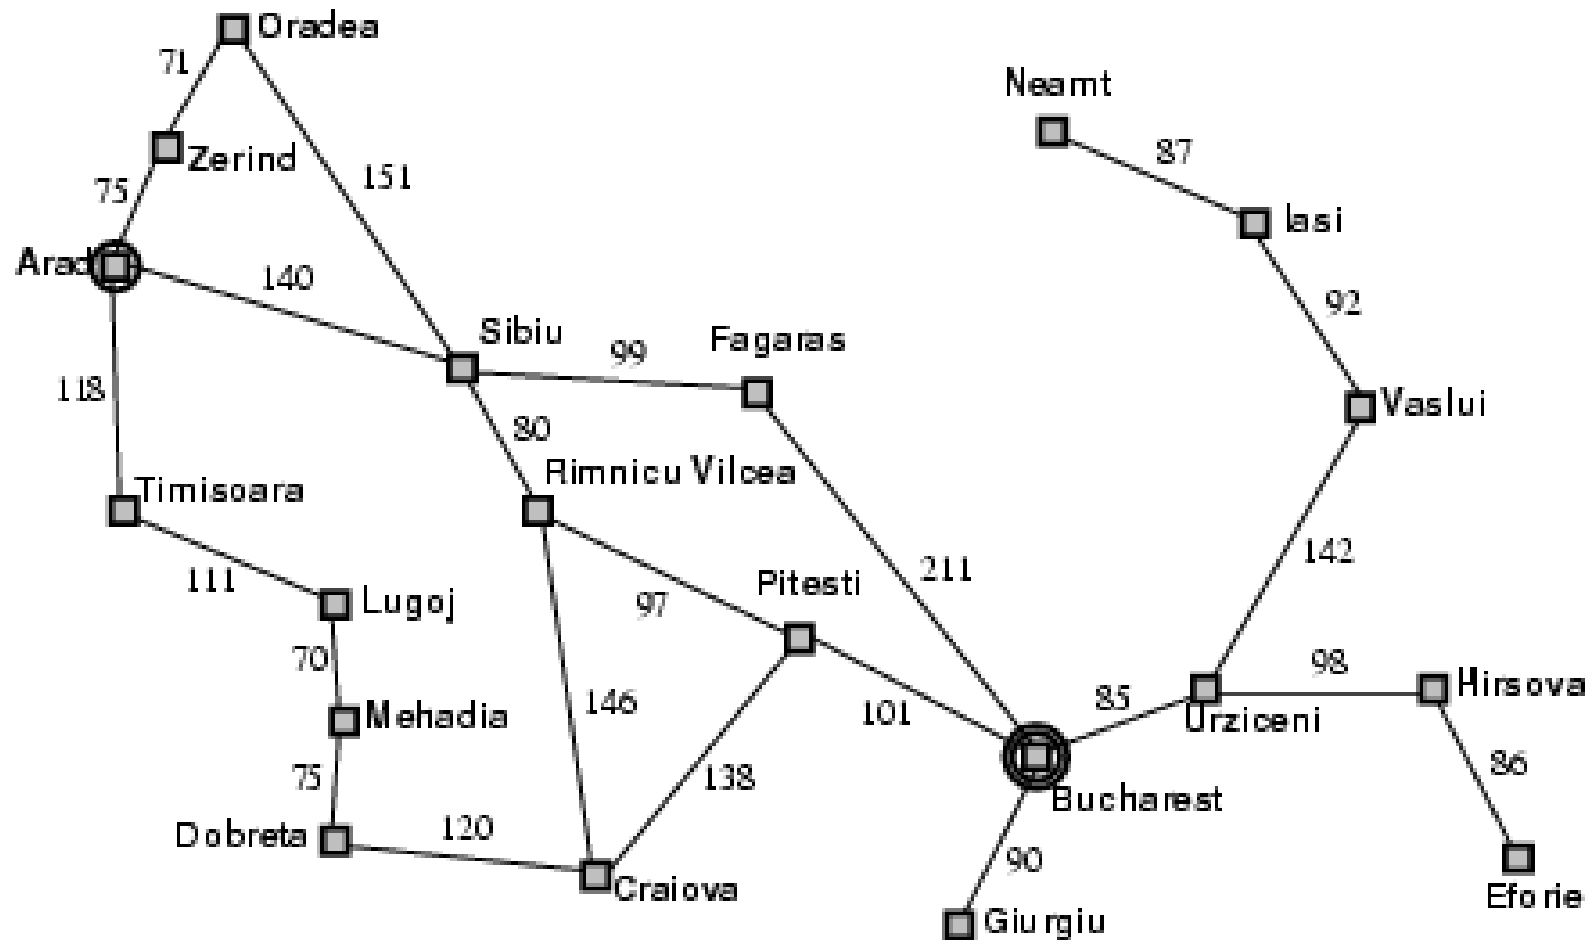
Best-First Search

- ❑ **Idea**: use an **evaluation function** $f(n)$ for each node
 - ❑ estimate of "desirability"
 - ❑ Expand most desirable unexpanded node
- ❑ **Implementation**: use a data structure that maintains the frontier in a decreasing order of desirability
- ❑ Is it really the best?
- ❑ **Special cases**: uniform-cost (Dijkstra's algorithm), greedy search, A* search
- ❑ A key component is a **heuristic function** $h(n)$:
 - ❑ $h(n)$ = estimated cost of the **cheapest path** from node n to a goal node
 - ❑ $h(n) = 0$ if n is the goal
 - ❑ $h(n)$ could be general or problem-specific

Best First Search Algorithm

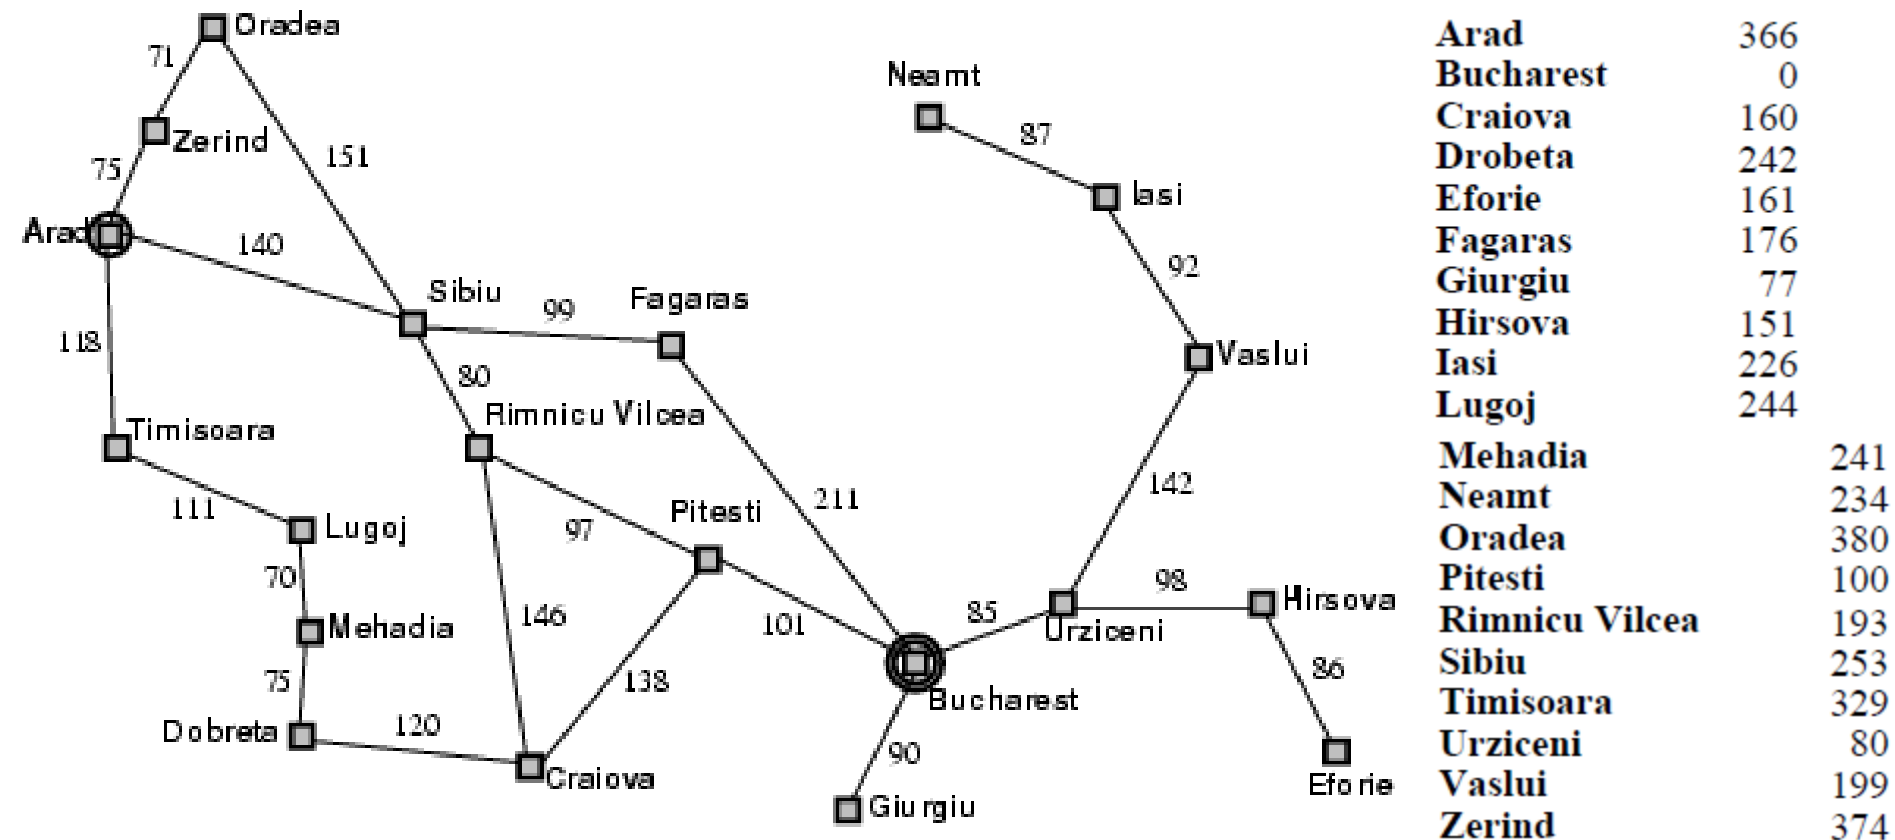
1. initialize the Q with the starting state (node)
2. while Q is not empty, do
 - 1) assign the first element of Q to N
 - 2) if N is the goal, return SUCCESS
 - 3) remove N from Q
 - 4) add the children of N to Q
 - 5) sort the entire Q by $f(n)$
3. return FAILURE

Recall Romania Map Example



What's a proper heuristic that measures cheapest path from current node to goal node?

Romania Map with Costs



Greedy Best-First Search

- ❑ Evaluation function: $f(n) = h(n)$
 - ❑ estimate the cost from n to goal
- ❑ h_{SLD} = straight line distance from n to Bucharest


Example: Arad to Bucharest



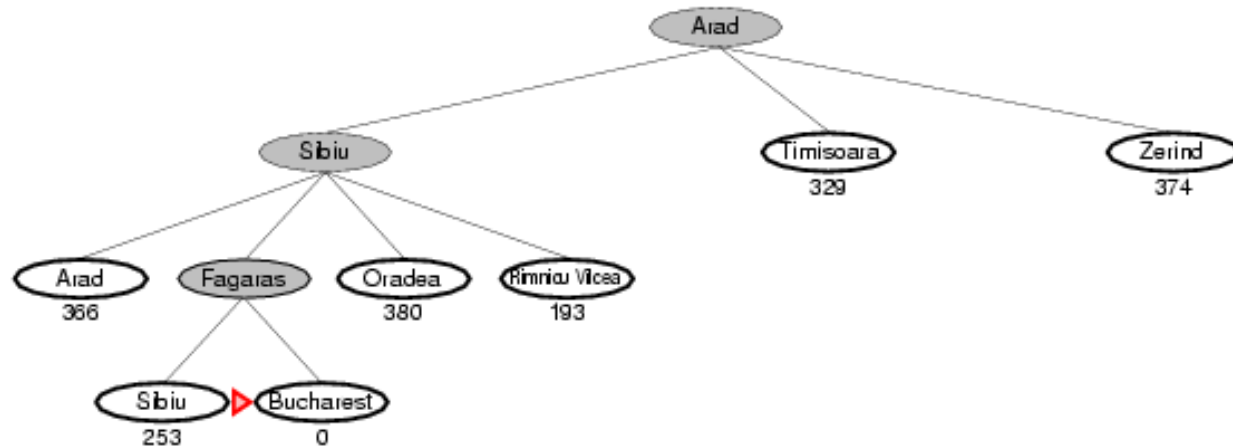
Example: Arad to Bucharest



Example: Arad to Bucharest



Example: Arad to Bucharest



Analysis of Greedy Best-First

❑ Complete?

- ❑ From Iasi to Fagaras
- ❑ No – can get stuck in loops, e.g., Iasi → Neamt → Iasi → Neamt → ...

❑ Time?

- ❑ $O(b^m)$, but a good heuristic can give dramatic improvement

❑ Space?

- ❑ $O(b^m)$ -- keeps all nodes in memory

❑ Optimal?

- ❑ No


In-Class Exercise #3.8

- Draw the search tree generated by Greedy search implemented with tree-search algorithm to find a path from T (Timisoara) to B (Bucharest)

A*: Minimizing Total Est. Cost

- ❑ Idea: avoid expanding paths that are already expensive
- ❑ **Evaluation function** $f(n) = g(n) + h(n)$
 - ❑ $g(n)$ = cost so far to reach n
 - ❑ $h(n)$ = estimated cost from n to goal
 - ❑ $f(n)$ = estimated total cost of path through n to goal


A* Search Example



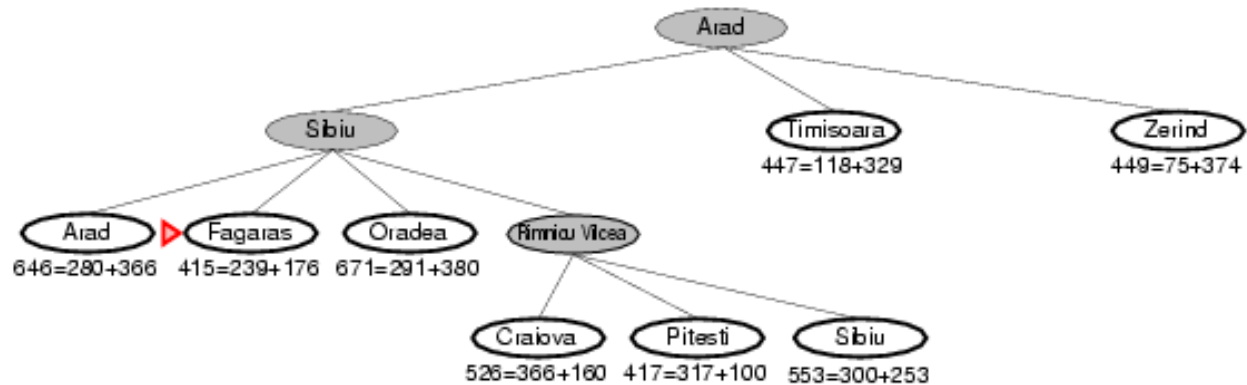
A* Search Example



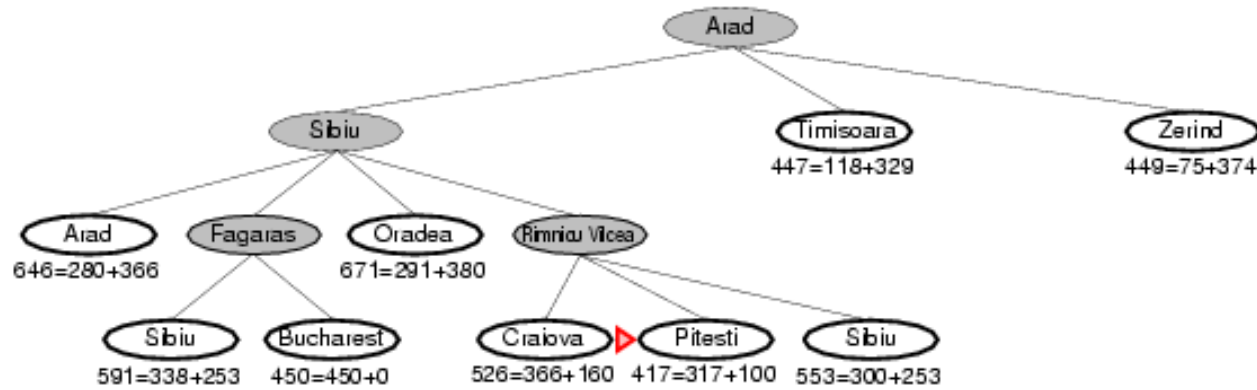
A* Search Example



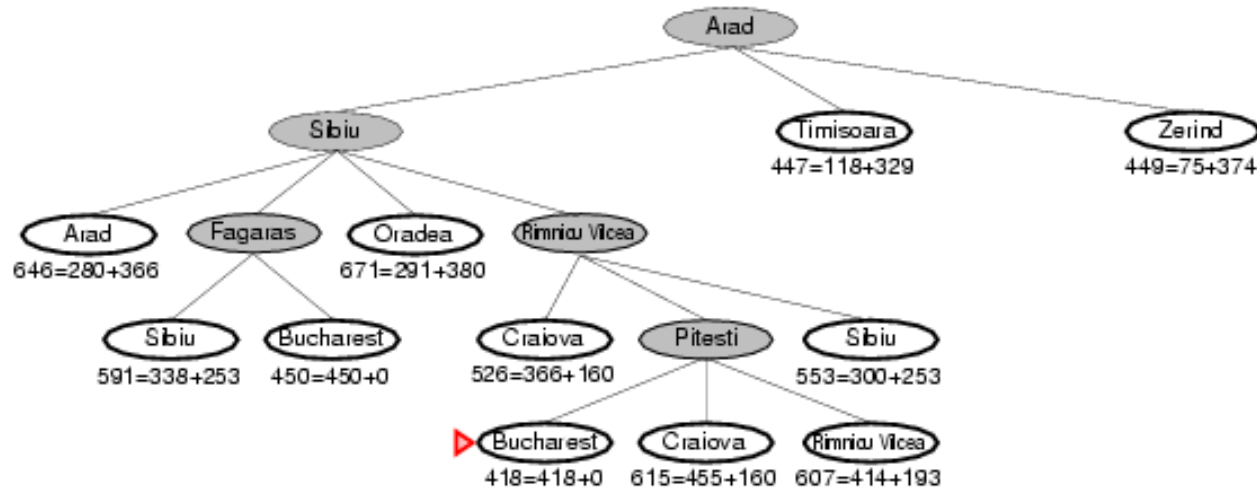
A* Search Example



A* Search Example



A* Search Example



In-Class Exercise #3.9


- Draw the search tree generated by applying A* and graph-search to find a path from Lugoj to Bucharest using the straight-line distance heuristic. Show the f-cost $f(n)$ for each node.

Admissible Heuristic

- A heuristic $h(n)$ is **admissible** if for every node n , $h(n) \leq h^*(n)$, where $h^*(n)$ is the **true** cost to reach the goal state from n
- An admissible heuristic **never overestimates** the cost to reach the goal, i.e., it is **optimistic**
- Example: $h_{SLD}(n)$ (never overestimates the actual road distance)
- **Theorem:** If $h(n)$ is admissible, A^* using TREE-SEARCH is optimal

Optimality of A* -- Proof

- Suppose some suboptimal goal G_2 has been generated and is in the frontier. Let n be an unexpanded node in the frontier such that n is on a shortest path to an optimal goal G .




- Assume the optimal cost is C^*
- $f(G_2) = g(G_2) + h(G_2) = g(G_2) > C^*$
- $f(n) = g(n) + h(n) \leq C^*$
- from the above, we have
 - $f(n) \leq C^* < f(G_2)$
- thus G_2 will not be expanded



Case-Study

- A* using Graph-Search returns a suboptimal solution with an $h(n)$ function that is admissible.



Consistency Heuristics

- A heuristic is **consistent** if for every node n , every successor n' of n generated by any action a , $h(n) \leq c(n,a,n') + h(n')$
- Triangle inequality
- Every consistent heuristic is also admissible


Proof by induction on the number k of nodes on the shortest path to any goal from n .

$K = 1$, let n' be the goal node; then $h(n) \leq c(n, a, n')$

Assume n' is on the shortest path k steps from the goal and that $h(n')$ is admissible by hypothesis, then:

$$h(n) \leq c(n,a,n') + h(n') \leq c(n,a,n') + h^*(n') = h^*(n)$$

So, $h(n)$ at $k+1$ steps from the goal is also admissible.




A* With Graph Search

- **Theorem:** If $h(n)$ is consistent, A* using GRAPH-SEARCH is optimal
- If h is consistent, and n' is a successor of n , we have
$$\begin{aligned}f(n') &= g(n') + h(n') \\&= g(n) + c(n, a, n') + h(n') \\&\geq g(n) + h(n) \\&= f(n)\end{aligned}$$
- i.e., $f(n)$ is non-decreasing along any path
- Whenever A* selects a node for expansion, the optimal path to that node has been found

Optimality of A*

- A* expands nodes in order of increasing f value
- Assume C^* is the optimal cost
 - A* expands all nodes with $f(n) < C^*$
 - A* might then expand some of the nodes right on the “goal contour” ($f(n) = C^*$) before selecting a goal node
 - Uniform-cost search ($h(n) = 0$) \rightarrow bands more circular



Analysis of A*

- ❑ Important idea:
 - ❑ appropriate $h(n)$ function
 - ❑ **A* is optimal efficient** (no other optimal alg. is guaranteed to expand fewer nodes than A*)
 - ❑ **pruning while still guaranteeing optimality**
- ❑ Complete?
 - ❑ Yes (unless there are infinitely many nodes with $f \leq f(G)$)
- ❑ Optimal?
 - ❑ Yes, with finite b and positive path cost
- ❑ However, A* is not the answer for all problems
- ❑ Time?
 - ❑ Exponential in the length of the solution
- ❑ Space?
 - ❑ Keeps all nodes in memory
 - ❑ A* usually runs out of space long before it runs out of time

Heuristic Functions

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- Two commonly used functions:
 - $h_1(n)$ = number of misplaced tiles
 - e.g., $h_1(n) = 8$ in the above example
 - $h_2(n)$ = sum of the distances of the tiles from their goal positions, called Manhattan distance
 - e.g., $h_2(n) = 3 + 1 + 2 + 2 + 2 + 3 + 3 + 2 = 18$ in the above example

Quality of Heuristic

- Assume # of nodes generated by A^* for a problem is N and the solution depth is d , then b^* is the effective branching factor that a uniform tree of depth d would have to have. Thus:
 - $N + 1 = 1 + b^* + (b^*)^2 + \dots + (b^*)^d$
- b^* might vary across problem instances, but is fairly constant for sufficiently hard problems
- A well-designed heuristic would have a value of b^* close to 1

The Comparison

	Search Cost			Effective Branching Factor		
d	IDS	$A^*(h_1)$	$A^*(h_2)$	IDS	$A^*(h_1)$	$A^*(h_2)$
2	10	6	6	2.45	1.79	1.79
4	112	13	12	2.87	1.48	1.45
6	680	20	18	2.73	1.34	1.30
8	6384	39	25	2.80	1.33	1.24
10	47127	93	39	2.79	1.38	1.22
12	364404	227	73	2.78	1.42	1.24
14	3473941	539	113	2.83	1.44	1.23
16	–	1301	211	–	1.45	1.25
18	–	3056	363	–	1.46	1.26
20	–	7276	676	–	1.47	1.27
22	–	18094	1219	–	1.48	1.28
24	–	39135	1641	–	1.48	1.26

Each data point corresponds to 100 instances of the 8-puzzle problem in which the solution depth varies.

Heuristics Domination

- If $h_2(n) \geq h_1(n)$ for all n (both admissible)
- then h_2 **dominates** h_1
- h_2 is better for search
 - A* using h_2 will never expand more nodes than A* using h_1

Inventing Admissible Heuristic

- ❑ A problem with fewer restrictions on the actions is called a **relaxed problem**
- ❑ The cost of an optimal solution to a relaxed problem is an admissible heuristic for the original problem
 - ❑ If the rules of the 8-puzzle are relaxed so that a tile can move **anywhere**, then $h_1(n)$ gives the shortest solution
 - ❑ If the rules are relaxed so that a tile can move to **any adjacent square**, then $h_2(n)$ gives the shortest solution

Inventing Admissible Heuristic

- If a collection of admissible heuristics $h_1 \dots h_m$ is available, and none of them dominates any of the others, we can choose
 - $h(n) = \max\{h_1(n), \dots, h_m(n)\}$
- We can also get from sub-problem of a given problem

In-Class Exercise #3.10

- The heuristic path algorithm is a best-first search in which the objective function is $f(n) = (2-w) \times g(n) + w \times h(n)$. For what values of w is the algorithm guaranteed to be optimal? (You may assume that h is admissible.) What kind of search does this perform when $w = 0$? When $w = 1$? When $w = 2$?

In-Class Exercise #3.11

- ❑ Prove each of the following statements:
 - ❑ Breadth-first search is a special case of uniform-cost search.
 - ❑ Breadth-first search, depth-first search, and uniform-cost search are special cases of best-first search.
 - ❑ Uniform-cost search is a special case of A^* search.

Discussion

- ❑ Project 1

Summary

- Applying heuristics to reduce search costs
 - Best-first search: $h(n)$
 - Greedy best-first search: $h(n)$
 - A*: $f(n) = g(n) + h(n)$