



Chapter 5: Adversarial Search

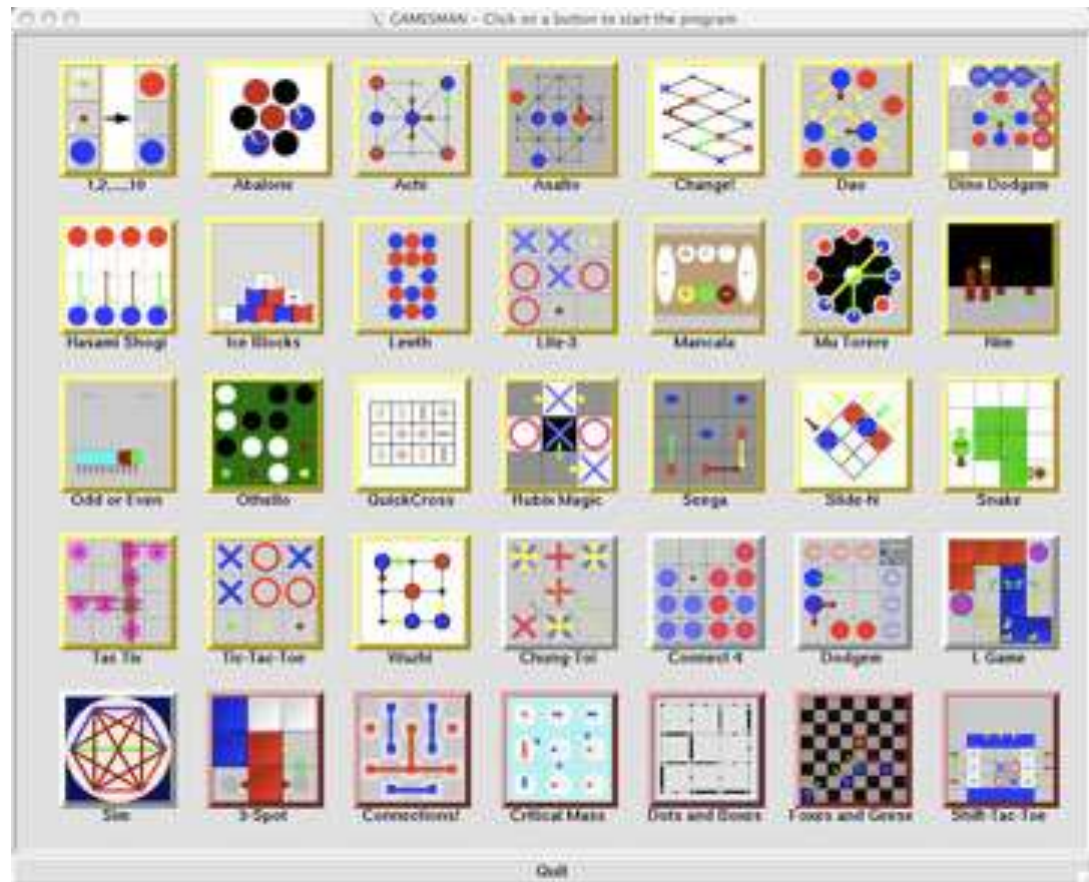
In which we examine the problems that arise when we try to plan ahead in a world where other agents are planning against us.

Adversarial Search

- **Multi-agent environment:**
 - any given agent needs to consider the actions of other agents and how they affect its own welfare
 - introduce possible contingencies into the agent's problem-solving process
 - cooperative vs. competitive
- **Adversarial search problems:** agents have conflicting goals -- games

Games vs. Search Problems

- "Unpredictable" opponent
 - specifying a move for every possible opponent reply
- Time limits
 - unlikely to find goal, must approximate



AI and Games

- In AI, “games” have special format:
 - deterministic, turn-taking, 2-player, zero-sum games of perfect information
 - Zero-sum describes a situation in which a participant’s gain or loss is exactly balanced by the losses or gains of the other participant(s)
 - Or, the total payoff to all players is the same for every instance of the game (constant sum)



Go! 围棋


Game Problem Formulation

- A game with 2 players (MAX and MIN, MAX moves first, turn-taking) can be defined as a search problem with:
 - **initial state**: board position
 - **player**: player to move
 - **successor function**: a list of legal (move, state) pairs
 - **goal test**: whether the game is over – terminal states
 - **utility function**: gives a numeric value for the terminal states (win, loss, draw)
- **Game tree = initial state + legal moves**



Optimal Strategies

- MAX must find a **contingent strategy**, specifying MAX's move in:
 - the initial state
 - the states resulting from every possible response by MIN
- E.g., **2-ply game** (the tree is one move deep, consisting of two half-moves, each of which is called a ply):



Minimax Value

- Perfect play for deterministic game, assume both players play optimally
- Idea: choose move to position with highest **minimax value** = best achievable payoff against best play

$$\text{MINIMAX-VALUE}(n) =$$

$$\text{Utility}(n)$$

if n is a terminal state


$$\max_{s \in \text{Successors}(n)} \text{MINIMAX}(s)$$

if n is a MAX node


$$\min_{s \in \text{Successors}(n)} \text{MINIMAX}(s)$$

if n is a MIN node

A Partial Game Tree for Tic-Tac-Toe



A Partial Game Tree for Tic-Tac-Toe



Minimax Algorithm

function MINIMAX-DECISION(*state*) *returns an action*

$v \leftarrow \text{MAX-VALUE}(\textit{state})$

return the *action* in SUCCESSORS(*state*) with value *v*

function MAX-VALUE(*state*) *returns a utility value*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow -\infty$

for *a, s* in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s))$

return *v*

function MIN-VALUE(*state*) *returns a utility value*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow \infty$

for *a, s* in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s))$

return *v*

Minimax with Tic-Tac-Toe

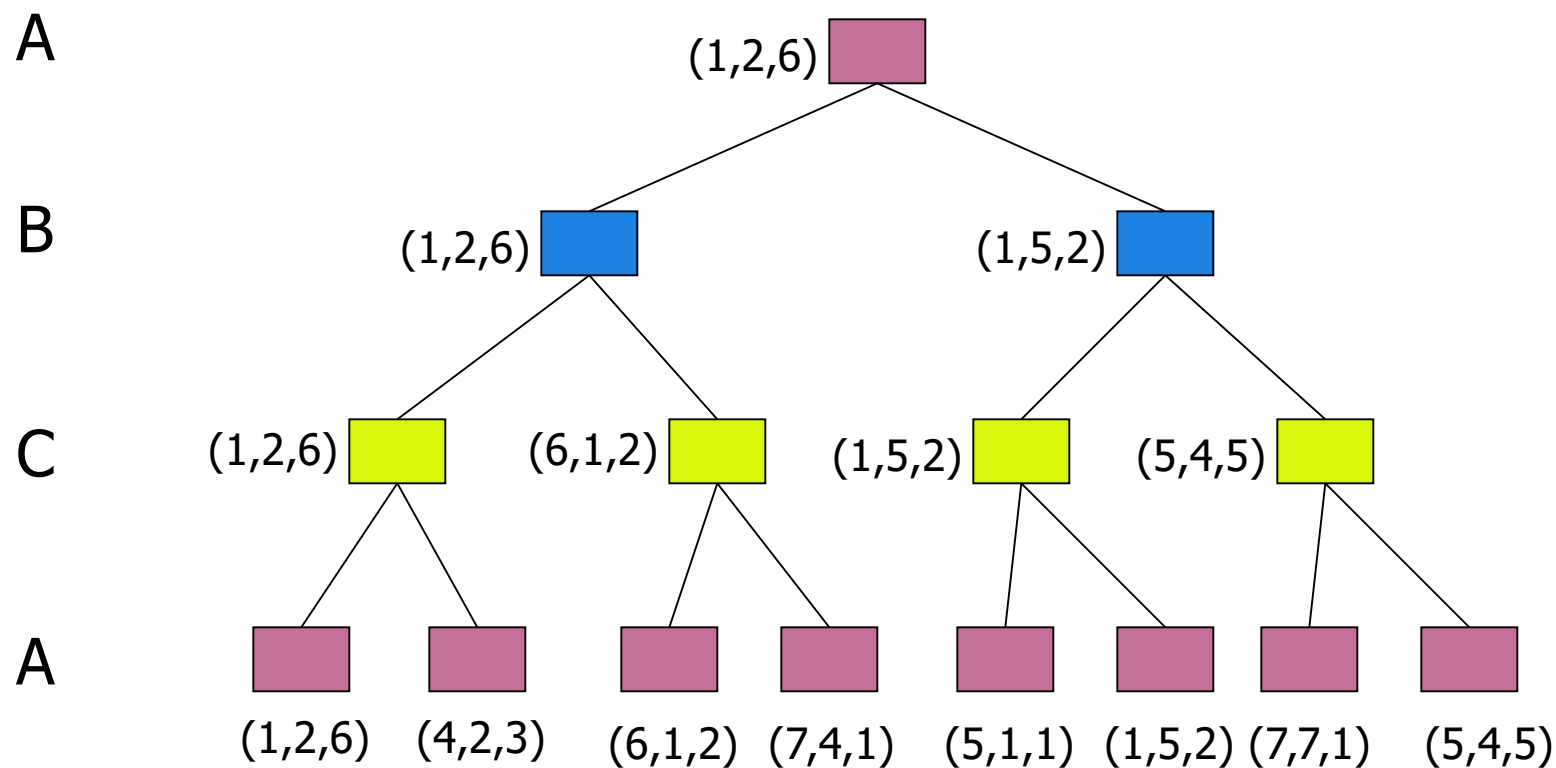
- Work on board.
- Walk through a real tic-tac-toe program.

Analysis of Minimax

- Optimal play for MAX assumes that MIN also plays optimally, what if MIN does not play optimally?
- A complete depth-first search?
 - Yes
- Time complexity?
 - $O(b^m)$
- Space complexity?
 - $O(bm)$ (depth-first exploration)
- For chess, $b \approx 35$, $m \approx 100$ for "reasonable" games
→ exact solution completely infeasible

Optimal Decisions for Multiplayer Games

- Extend minimax idea to multiplayer





Interesting Thoughts

- Multiplayer games usually involve alliances, which can be a natural consequence of optimal strategies
- If the game is non zero-sum, collaboration can also occur
 - For example, a terminal state with utilities $\langle V_a = 1000, V_b = 1000 \rangle$
 - The optimal strategy is for both players to do everything possible to reach this state



α - β Pruning

- The number of game states with minimax search is exponential in the # of moves
- Is it possible to compute the correct minimax decision without looking at every node in the game tree?
- Need to **prune** away branches that cannot possibly influence the final decision



α - β Pruning Example





α - β Pruning Example





α - β Pruning Example





α - β Pruning Example





α - β Pruning Example



α - β Pruning Example



α - β Pruning Example



A Simple Formula

$$\text{MINIMAX} - \text{VALUE}(\text{root})$$

$$= \max(\min(3, 12, 8), \min(2, x, y), \min(14, 5, 2))$$


$$= \max(3, \min(2, x, y), 2)$$

$$= \max(3, z, 2) \quad \text{where } z \leq 2$$

$$= 3$$

The value of the root and hence the minimax decision are independent of the values of pruned leaves x and y

α - β Pruning Exercise



Basic Idea of α - β

- Consider a node **n** such that Player has a choice of moving to
- If Player has a **better choice m** either at the parent of n or at any choice point further up, then **n will never be reached in actual play**
- α - β pruning gets its name from the two parameters that describe bounds on the backed-up values


Player

Opponent

...

Player

Opponent



Definitions of α and β

- α = the value of the best (highest-value) choice we have found so far at any choice point along the path for MAX
- β = the value of the best (lowest-value) choice we have found so far at any choice point along the path for MIN
- α - β search updates the values of α and β as it goes along and prunes the remaining branches at a node as soon as the value of the current node is worse than the current α or β for MAX or MIN respectively

function ALPHA-BETA-SEARCH(*state*) *returns an action*

inputs: *state*, current state in game

$v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$

return the *action* in SUCCESSORS(*state*) with value *v*

function MAX-VALUE(*state*, α , β) *returns a utility value*

inputs: *state*, current state in game

α , the value of the best alternative for MAX along the path to *state*

β , the value of the best alternative for MIN along the path to *state*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow -\infty$

for *a, s* in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s, \alpha, \beta))$

if $v \geq \beta$ **then return** *v*

$\alpha \leftarrow \text{MAX}(\alpha, v)$

return *v*

function MIN-VALUE(*state*, α , β) *returns a utility value*

inputs: *state*, current state in game

α , the value of the best alternative for MAX along the path to *state*

β , the value of the best alternative for MIN along the path to *state*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow +\infty$

for *a, s* in SUCCESSORS(*state*) **do**


$v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s, \alpha, \beta))$

if $v \leq \alpha$ **then return** *v*


$\beta \leftarrow \text{MIN}(\beta, v)$

return *v*

Now, Trace The Behavior



Analysis of α - β



- Pruning **does not** affect final result
- The effectiveness of alpha-beta pruning is highly dependent on the **order of successors**
- It might be worthwhile to try to examine first the successors that are likely to be best
- With "perfect ordering," time complexity = $O(b^{m/2})$
 - **effective branching factor** becomes \sqrt{b}
 - For chess, 6 instead of 35
 - it can look ahead roughly twice as far as minimax in the same amount of time
 - Ordering in chess: **captures, threats, forward moves, and then backward moves**

Random Ordering?

- If successors are examined in random order rather than best-first, the complexity will be roughly $O(b^{3m/4})$
- Adding dynamic move-ordering schemes, such as trying first the moves that were found to be best last time, brings us close to the theoretical limit
- The best moves are often called **killer moves** (killer move heuristic)

Dealing with Repeated States

- In games, repeated states occur frequently because of **transpositions** --
 - different permutations of the move sequence end up in the same position
 - e.g., [a1, b1, a2, b2] vs. [a1, b2, a2, b1]
- It's worthwhile to store the evaluation of this position in a hash table the first time it is encountered
 - similar to the "**explored set**" in graph-search
- Tradeoff:
 - Transposition table can be too big
 - Which to keep and which to discard

Imperfect, Real-Time Decisions

- Minimax generates the entire game search space
- Alpha-beta prunes large part of it, but still needs to search all the way to terminal states
- However, moves must be made in reasonable amount of time
- **Standard approach**: turning non-terminal nodes into terminal leaves
 - **cutoff test**: replaces terminal test, e.g., depth limit
 - **heuristic evaluation function** = estimated desirability or utility of position

Evaluation Functions

- The performance of a game-playing program is dependent on the **quality of its evaluation function**
 - Order the terminal states the same way as the true utility function
 - Evaluation of nonterminal states correlate with the actual chance of winning
 - Computation must not take too long!

Playing Chess, Experience is Important

- Most eval. Functions work by calculating various features of the state
- Different features form various categories of states
- Consider the category of two-pawns vs. one pawn and suppose experience suggests that 72% of the states lead to a win (+1), 20% to a loss(0), and 8% to a draw (1/2)
 - Expected value = $0.72*1 + 0.20*0 + 0.08*1/2 = 0.76$
- Too many categories and too much experience are required

-
- Introductory chess books give an approximate material value for each piece:
 - each pawn is worth 1, a knight or bishop is worth 3, a rook 5, and the queen 9
 - Other features such as “good pawn structure” and “king safety” worth half a pawn
 - For chess, typically **linear** weighted sum of **features**
$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$
 - w_i is a weight and f_i is a feature of the position
 - e.g., f is the # of each piece, and w is the value of each piece
 - Linear assumption: contribution of each feature is independent of the values of the other features
 - There are also non-linear combinations of features

Cutting off Search

- Modify alpha-beta search so that:
 - Terminal? is replaced by Cutoff?
 - Utility is replaced by Eval
 - **if Cutoff-Test**(state, depth) **then return Eval**(state)
 - depth is chosen such that the amount of time used will not exceed what the rules of the game allow
 - Iterative deepening search can be applied
 - When time runs out, returns the move selected by the deepest completed search

Result?

- Does it work in practice?
 - Chess problem. Assume we can generate and evaluate a million nodes per second, this will allow us to search roughly 200 million nodes per move under standard time control (3 minutes per move).
 - With minimax, only 5-ply, but with alpha-beta, 10-ply
- 4-ply lookahead is a hopeless chess player!
 - 4-ply \approx human novice
 - 8-ply \approx typical PC, human master
 - 12-ply \approx Deep Blue, Kasparov

Deterministic Games in Practice

- **Checkers**: Chinook ended 40-year-reign of human world champion Marion Tinsley in 1994. Used a precomputed endgame database defining perfect play for all positions involving 8 or fewer pieces on the board, a total of 444 billion positions.
- **Chess**: Deep Blue defeated human world champion Garry Kasparov in a six-game match in 1997. Deep Blue searches 200 million positions per second, uses very sophisticated evaluation, and undisclosed methods for extending some lines of search up to 40 ply.
- **Othello**: human champions refuse to compete against computers, who are too good.
- **Go**: human champions refuse to compete against computers, who are too bad. In go, $b > 300$, so most programs use pattern knowledge bases to suggest plausible moves.

Summary

- Games are fun to work on!
- They illustrate several important points about AI
- Perfection is unattainable → must approximate

In-Class Exercise #5.1

- Prove that: for every game tree, the utility obtained by MAX using minimax decisions against a suboptimal MIN will never be lower than the utility obtained playing against an optimal MIN.

In-Class Exercise 5.2

- Consider the following game:
 - Draw the complete game tree, using the following convention:
 - State: (S_a, S_b) where S_a and S_b denote the token locations
 - Identify the terminal state in a square box and assign it a value
 - Put loop states in double square boxes
 - Since it's not clear how to assign values to loop states, annotate each with a "?"

